

Dynamic Memory based Attention Network for Sequential Recommendation

**Qiaoyu Tan ¹, Jianwei Zhang ², Ninghao Liu ¹, Xiao Huang ³
Hongxia Yang ², Jingren Zhou ², Xia Hu ¹**

¹ Texas A&M University, ² Alibaba Group, ³ The Hong Kong Polytechnic University
{qytan,nhliu43,xiahu}@tamu.edu, xiaohuang@comp.polyu.edu.hk
{zhangjianwei.zjw, yang.yhx, jingren.zhou}@alibaba-inc.com

Problem Formulation

- **Sequential Recommendation:** Assume \mathcal{U} and \mathcal{V} denote the sets of users and items, respectively. $S = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{|S|}\}$ represents the behavior sequence in chronological order of a user. $\mathbf{x}_t \in \mathcal{V}$ records the t -th item interacted by the user. Given an observed behavior sequence $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{|S|}\}$, the sequential recommendation task is to predict the next items that the user might be interacted with.

Challenges

- Given that the response time in real-world systems is limited, it has become expensive to scan over the entire behavior sequence at each prediction time.
- It is crucial to model the whole behavior sequence for a more accurate recommendation.
- It is hard to explicitly control the contributions of long-term or short-term interests for user modeling.

Methodology

- Truncating the whole user behavior sequence into several successive sub-sequences and optimizes the model sequence by sequence.
- A recurrent attention network and a long-term attention network with dynamic memory are derived for short-term and long-term interest modeling.
- A gate mechanism is applied to adaptively control the importance of the above two interests combination.

Dynamic memory-based self-attention network (DMAN)

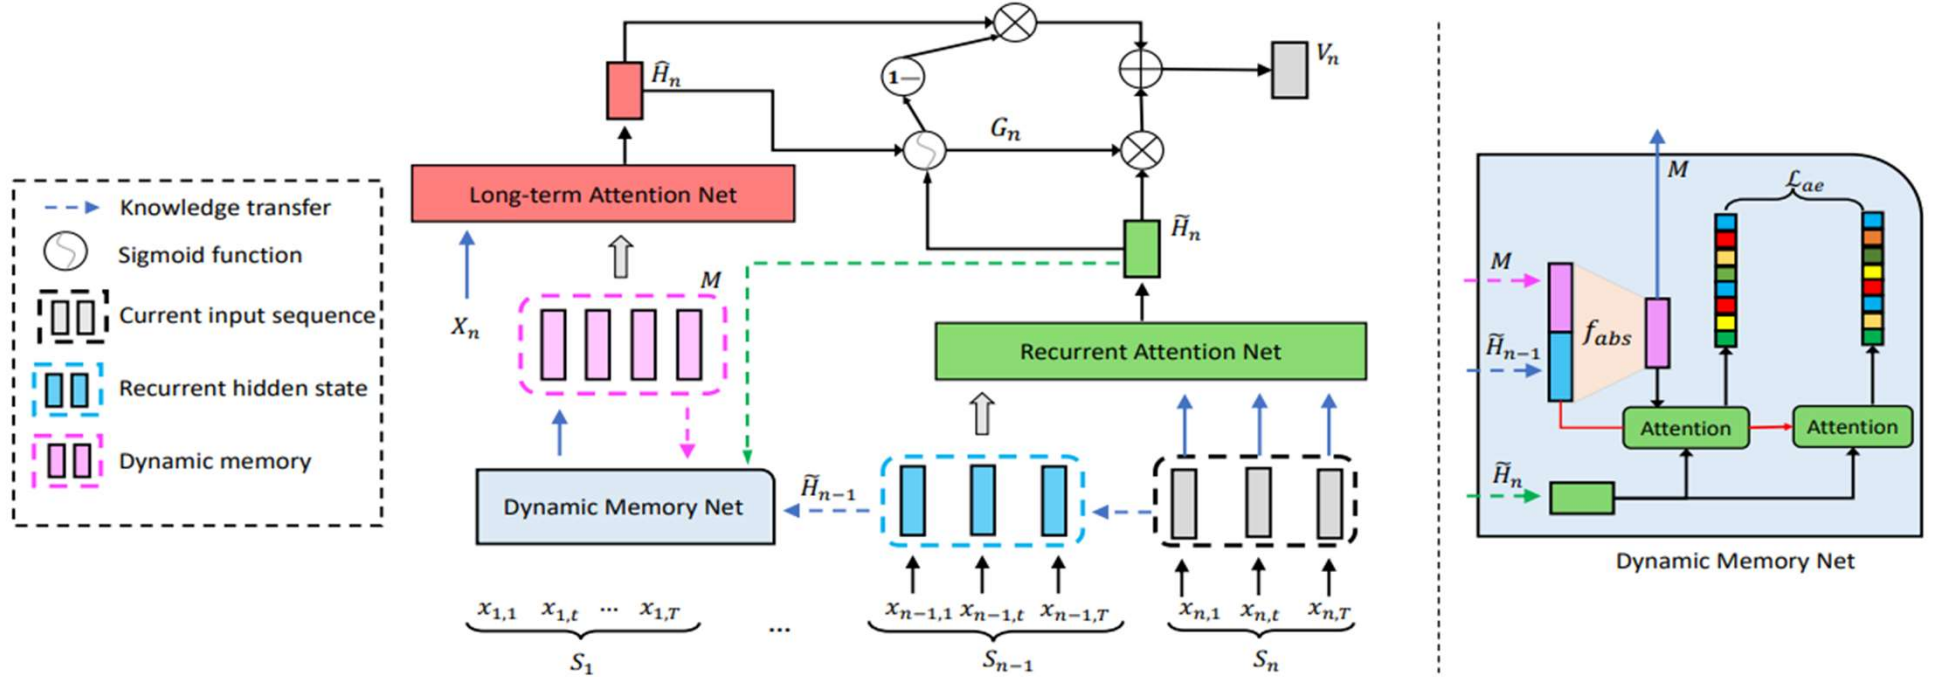


Figure 1: Illustration of DMAN for one layer. It takes a series of sequences as input and trains the model sequence by sequence. When processing the n -th sequence S_n , the recurrent attention network is applied to extract short-term user interest by using the previous hidden state \tilde{H}_{n-1} as context. Meanwhile, the long-term attention network is utilized to extract long-term interest based on the memory blocks M . Next, the short-term and long-term interests are combined via a neural gating network for joint user modeling. Finally, the dynamic memory network updates the memory blocks via fusing the information in \tilde{H}_{n-1} , and the model continues to process the next sequence. The overall model is optimized by maximizing the likelihood of observed sequence, while the dynamic memory network is trained based on a local reconstruction loss \mathcal{L}_{ae} .

Recurrent Attention Network

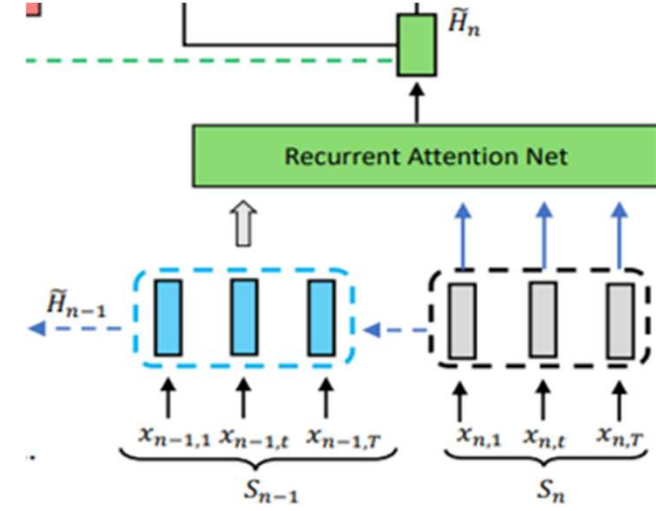
Truncating accumulated behavior sequence S into a series of successive sub-sequences with fixed window size T , i.e., $S = \{S_n\}_{n=1}^N$. Thus, $S_n = \{\mathbf{x}_{n,1}, \mathbf{x}_{n,2}, \dots, \mathbf{x}_{n,T}\}$ denotes the n -th sequence. Let $\tilde{\mathbf{H}}_{n-1}^l \in \mathbb{R}^{T \times D}$ denote the l -th layer hidden state produced for sequence S_{n-1} . S_n can be calculated as follows.

$$\tilde{\mathbf{H}}_n^l = \text{Atten}_{rec}^l(\tilde{\mathbf{Q}}_n^l, \tilde{\mathbf{K}}_n^l, \tilde{\mathbf{V}}_n^l) = \text{softmax}\left(\tilde{\mathbf{Q}}_n^l (\tilde{\mathbf{K}}_n^l)^T\right) \tilde{\mathbf{V}}_n^l$$

$$\tilde{\mathbf{Q}}_n^l = \tilde{\mathbf{H}}_n^{l-1} \tilde{\mathbf{W}}_Q^T, \quad \tilde{\mathbf{K}}_n^l = \mathbf{H}_n^{l-1} \tilde{\mathbf{W}}_K^T, \quad \tilde{\mathbf{V}}_n^l = \mathbf{H}_n^{l-1} \tilde{\mathbf{W}}_V^T$$

$$\mathbf{H}_n^{l-1} = \tilde{\mathbf{H}}_n^{l-1} || \text{SG}(\tilde{\mathbf{H}}_{n-1}^{l-1})$$

The function $\text{SG}(\cdot)$ stands for stop-gradient from previous hidden state $\tilde{\mathbf{H}}_{n-1}^{l-1}$. The input of the first layer is the sequence embedding matrix $\mathbf{X}_n = \{\mathbf{x}_{n,1}, \mathbf{x}_{n,2}, \dots, \mathbf{x}_{n,T}\} \in \mathbb{R}^{T \times D}$. The final short-term interest embedding is defined as $\tilde{\mathbf{H}}_n = \tilde{\mathbf{H}}_n^L$.




Long-term Attention Network

Define the l -th layer memory matrix $\mathbf{M}^l \in \mathbb{R}^{T \times D}$ to explicitly memorize a user's long-term preferences, where m is the number of memory slots. The long-term hidden state $\hat{\mathbf{H}}_n^l$ can be estimated as follows.

$$\hat{\mathbf{H}}_n^l = \text{Atten}^l(\hat{\mathbf{Q}}_n^l, \hat{\mathbf{K}}_n^l, \hat{\mathbf{V}}_n^l) = \text{softmax}\left(\hat{\mathbf{Q}}_n^l, (\hat{\mathbf{K}}_n^l)^T\right) \hat{\mathbf{V}}_n^l$$

$$\hat{\mathbf{Q}}_n^l = \hat{\mathbf{H}}_n^{l-1} \hat{\mathbf{W}}_Q^T, \quad \hat{\mathbf{K}}_n^l = \mathbf{M}^{l-1} \hat{\mathbf{W}}_K^T, \quad \hat{\mathbf{V}}_n^l = \mathbf{M}^{l-1} \hat{\mathbf{W}}_V^T$$

The input of the first layer of query is \mathbf{X}_n . The final long-term interest embedding is denoted as $\hat{\mathbf{H}}_n = \hat{\mathbf{H}}_n^L$.



Dynamic Memory Network

To abstract a user's long-term interests from the past actively, the memory is updated by an abstraction function

$$f_{abs}: \mathbb{R}^{(m+T) \times D} \rightarrow \mathbb{R}^{m \times D},$$

$$\mathbf{M}^l \leftarrow f_{abs}(\mathbf{M}^l, \tilde{\mathbf{H}}_{n-1}^l)$$

- **Abstraction function f_{abs}**

f_{abs} is parameterized with the dynamic routing model in CapsNet for user's divers interest capture. Referring capsules from the first layer and second layer as primary capsules (input vector $\mathbf{x}_i, i \in \{1, \dots, T + m\}$) and interest capsules (output vector $\bar{\mathbf{x}}_j, j \in \{1, \dots, m\}$). The dynamic routing logit b_{ij} between primary capsule i and interest capsule j is computed by

$$b_{ij} = \bar{\mathbf{x}}_j^T \mathbf{W}_{ij} \mathbf{x}_i, \quad \bar{\mathbf{x}}_j = \text{squash}(\mathbf{s}_j) = \frac{\|\mathbf{s}_j\|^2}{1 + \|\mathbf{s}_j\|^2} \frac{\mathbf{s}_j}{\|\mathbf{s}_j\|}$$

$$\mathbf{s}_j = \sum_{i=1}^{m+T} \alpha_{ij} \mathbf{W}_{ij} \mathbf{x}_i, \quad \alpha_{ij} = \exp(b_{ij}) / \sum_{j'=1}^{m+T} \exp(b_{ij'})$$

The routing process usually repeats three times to converge. When routing finishes, the output interest capsules of user u are then used as the memory, i.e., $\mathbf{M} = [\bar{\mathbf{x}}_1, \dots, \bar{\mathbf{x}}_m]$.

Dynamic Memory Network


- **Attention-based reconstruction loss \mathcal{L}_{ae}**

To make the primary interests can be extracted by f_{abs} as much as possible, an auxiliary attention-based reconstruction loss \mathcal{L}_{ae} is designed as follows.


$$\mathcal{L}_{ae} = \min \sum_{l=1}^L ||\text{atten}_{rec}^l(\tilde{\mathbf{Q}}^l, \tilde{\mathbf{K}}^l, \tilde{\mathbf{V}}^l) - \text{atten}_{rec}^l(\tilde{\mathbf{Q}}^l, \hat{\mathbf{K}}^l, \hat{\mathbf{V}}^l)||_F^2$$

$$\tilde{\mathbf{Q}}^l = \tilde{\mathbf{H}}_n^l, \quad \tilde{\mathbf{K}}^l = \tilde{\mathbf{V}}^l = \mathbf{M}_{pre}^l || \tilde{\mathbf{H}}_{n-1}^l, \quad \hat{\mathbf{K}}^l = \hat{\mathbf{V}}^l = \mathbf{M}_{up}^l$$

The recurrent attention network is reused here but keep the parameters fixed and not trainable.



Neural Gating Network



Considering that a user's future intention can be influenced by early behaviors, while short-term and long-term interests may contribute differently for next-item prediction over time, a neural gating network is applied to adaptively control the importance of the two interest embeddings.

$$\mathbf{V}_n = \mathbf{G}_n \odot \tilde{\mathbf{H}}_n + (1 - \mathbf{G}_n) \odot \hat{\mathbf{H}}_n$$

$$\mathbf{G}_n = \sigma(\tilde{\mathbf{H}}_n \mathbf{W}_{short} + \hat{\mathbf{H}}_n \mathbf{W}_{long})$$

The final user embedding $\mathbf{V}_n \in \mathbb{R}^{T \times D}$ is obtained by a feature-level weighted sum of two types of interest embeddings controlled by the gate.

Model Optimization

Given the training sample (u, t) in a sequence S_n with the user embedding vector $\mathbf{V}_{n,t}$ and target item embedding \mathbf{x}_t , the model aims to minimize the following negative likelihood

$$\mathcal{L}_{like} = - \sum_{u \in \mathcal{U}} \sum_{t \in S_n} \log P(\mathbf{x}_{n,t} | \mathbf{x}_{n,1}, \mathbf{x}_{n,2}, \dots, \mathbf{x}_{n,t-1}) = - \sum_{u \in \mathcal{U}} \sum_{t \in S_n} \log \frac{\exp(\mathbf{x}_t^T \mathbf{V}_{n,t})}{\sum_{j \in \mathcal{V}} \exp(\mathbf{x}_j^T \mathbf{V}_{n,t})}$$

A negative sampling strategy is adopted to approximate the softmax function in experiments. The \mathcal{L}_{ae} and \mathcal{L}_{like} are separately updated in order to preserve long-term interests better. Specifically, the \mathcal{L}_{like} is updated first by feeding a new sequence and then updating the abstraction function's parameters by minimizing \mathcal{L}_{ae} .

Experiments

Table 3: Sequential recommendation performance over three benchmarks. * indicates the model only use the latest behavior sequence for training; otherwise, the whole behavior sequence. The second best results are underlined.

| Models | MovieLens | | | Taobao | | | JD.com | | |
|----------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| | HR@10 | HR@50 | NDCG@100 | HR@50 | HR@100 | NDCG@100 | HR@10 | HR@50 | NDCG@100 |
| GRU4Rec* | 17.69 | 43.13 | 16.90 | 10.42 | 14.01 | 4.23 | 27.65 | 38.73 | 23.40 |
| Caser* | 18.98 | 45.64 | 17.62 | 13.71 | 16.51 | 6.89 | 29.27 | 40.16 | 24.25 |
| SASRec* | 21.02 | 47.28 | 19.05 | 16.41 | 22.83 | 9.23 | 33.98 | 44.89 | 27.41 |
| GRU4Rec | 19.78 | 47.40 | 18.75 | 13.48 | 16.53 | 5.81 | 35.28 | 47.52 | 27.64 |
| Caser | 20.80 | 48.12 | 19.28 | 15.55 | 17.91 | 7.35 | 36.76 | 49.13 | 28.35 |
| SASRec | 22.96 | 50.09 | 20.36 | 20.47 | 24.48 | 9.84 | 38.99 | 52.64 | 31.32 |
| SHAN | 21.34 | 49.52 | 19.55 | 18.87 | 21.94 | 8.73 | 37.72 | 50.55 | 29.80 |
| HPMN | 22.84 | 50.54 | 19.77 | 19.98 | 24.37 | 9.66 | 39.14 | 53.22 | 32.24 |
| SDM | <u>23.42</u> | <u>51.26</u> | <u>20.44</u> | <u>21.66</u> | <u>25.42</u> | <u>10.22</u> | <u>40.68</u> | <u>55.30</u> | <u>34.82</u> |
| DMAN | 25.18 | 53.24 | 22.03 | 24.92 | 29.37 | 11.13 | 44.58 | 58.82 | 36.93 |
| Improv. | 7.51% | 3.86% | 7.77% | 15.05% | 15.53% | 8.90% | 9.58% | 6.36% | 6.05% |


Table 4: Performance on long user behavior data XLong.

| Method | Recall@200 | Recall@500 |
|----------|--------------|--------------|
| GRU4Rec* | 0.079 | 0.098 |
| Caser* | 0.084 | 0.105 |
| SASRec* | 0.105 | 0.123 |
| GRU4Rec | 0.046 | 0.063 |
| Caser | 0.023 | 0.041 |
| SASRec | 0.061 | 0.096 |
| SHAN | 0.091 | 0.115 |
| HPMN | 0.118 | 0.136 |
| SDM | 0.107 | 0.129 |
| DMAN | 0.132 | 0.163 |

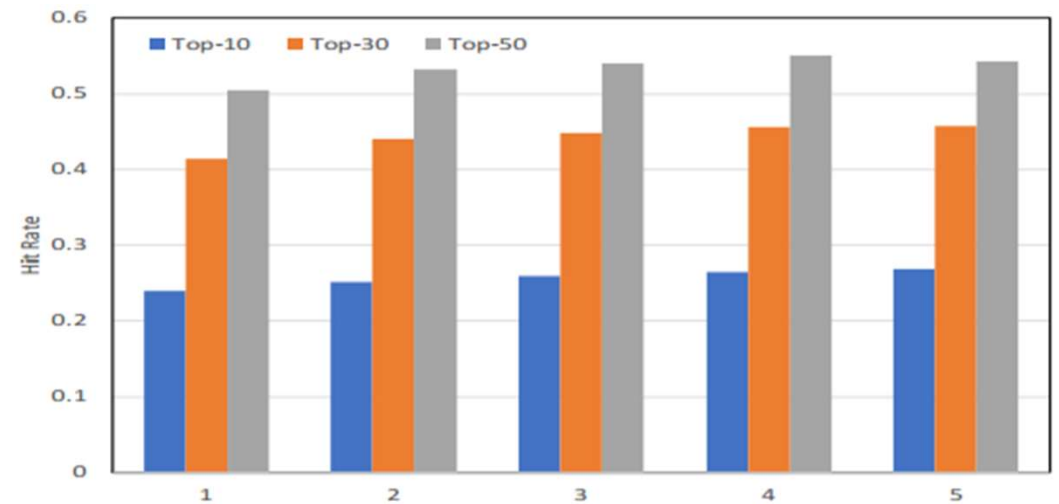
Table 5: Ablation study of DMAN.

| Dataset | Method | Recall@100 | NDCG@100 |
|---------|-----------|------------|----------|
| Taobao | DMAN-XL | 0.237 | 0.094 |
| | DMAN-FIFO | 0.263 | 0.108 |
| | DMAN-NRMA | 0.257 | 0.104 |
| | DMAN | 0.293 | 0.111 |
| XLong | DMAN-XL | 0.021 | 0.013 |
| | DMAN-FIFO | 0.036 | 0.017 |
| | DMAN-NRAN | 0.043 | 0.019 |
| | DMAN | 0.054 | 0.022 |

Experiments



(a) Memory slots m



(b) Layer size L